

Kubernetes Advanced



kubernetes

Monitoring




kubernetes

Session Contents



- Use Kubectl
- Kubernetes Dashboard
- Log management
- Prometheus & Grafana
- Autoscaling

Use Kubectl



Kubectl for monitoring



- Many times kubectl can be the first (and only) tool available to access the cluster
- Cannot be used to have a proactive approach to monitoring but can be used to get detailed understanding about cluster issues
- With the needed permissions you can have complete understanding about cluster behavior

Kubectl describe



```
kubectl describe pod <pod> [-n <namespace>]
```

- Shows details about pod
 - Metadata
 - Network
- Lists all events occurred during pod lifecycle
- First place to go when pod don't have "Running" status

Kubectl logs



```
kubectl logs <pod> [-n <namespace>]
```

- Shows pod stdout and stderr
- Flag `-f` blocks the console and show new lines

Kubectl port-forward



```
kubectl port-forward pod <pod> [-n <ns>] hostport:podPort  
kubectl port-forward svc <svc> [-n <ns>] hostport:podPort
```

- Maps a port on machine with pod port
- Allow to make direct requests
- When using service, maps directly to only one container (no load balancing)

Kubectl top



```
kubectl top node <node>
```

```
kubectl top pod <pod> [-n <ns>]
```

- Display resource (CPU/memory) usage of the resources (nodes or pods)
- Due to the metrics service delay, they may be unavailable for a few minutes since pod creation
- Use native Kubernetes metrics server

Kubernetes Dashboard



Kubernetes Dashboard



- Web-based Kubernetes user interface to have a more user-friendly way to look into your cluster.
- Kubernetes Dashboard were created and is maintained by Kubernetes Community
- Initially, was the only Web-based tool to monitor your cluster
- Now, is not so used on production environment
 - Newer and better tools arrive on Kubernetes Landscape
 - Limitation on metrics since it uses only Kubernetes Vanilla metrics
- You can use Dashboard to deploy containerized applications to a Kubernetes cluster, troubleshoot your containerized application, and manage the cluster resources.

Kubernetes Dashboard




- You can use Dashboard to:
 - Monitor your cluster resources
 - Manage Kubernetes resources
 - Get an overview of applications running on your cluster
 - Troubleshoot your containerized application
 - Deploy containerized applications to a Kubernetes cluster
- Provide wizards to scale a Deployment, initiate a rolling update, restart a pod or deploy new applications using a deploy wizard.

Demo | Kubernetes Dashboard



Log Management



Motivation



- Containers runs on top of an ephemeral layer that is deleted each time a pod is deleted
- If you write your logs to a file in this layer you may lose them
- Even during execution can be hard to reach them
- How to have access to these logs and keep them for as long as needed?






















Logs on Kubernetes



- As a best practice, everything that needs to be logged should be write to standard output or standard error of each container
- Log Management tools for Kubernetes uses a concept of creating a DaemonSet to have a pod in each node to access to those streams
- Then, after collecting the data, send them to a centralized server where the logs are kept for as long as needed
- Finally, the full solution have a visualization layer where logs can be queried and accessed from outside of the cluster

CNCF Logging



 Alibaba Cloud Log Service Alibaba Cloud Log Service Alibaba Cloud MCap: \$214.1B	 DataSet Scalyr Funding: \$27.6M	 elastic Elastic ★ 62,985 MCap: \$5.1B	 fluentd Cloud Native Computing Foundation (CNCF) ★ 11,831 Funding: \$3M	 Grafana loki Grafana Labs ★ 18,533 Funding: \$535.2M	 graylog Graylog ★ 6,450 Funding: \$27.4M
 humio A CrowdStrike Company Humio Funding: \$31.8M	 Loggie NetEase ★ 887 MCap: \$53.4B	 LOGGLY Loggly Funding: \$47.4M	 LOGIQ Logiq.ai Funding: \$1.8M	 logstash Elastic ★ 13,322 MCap: \$5.1B	 mezmo Mezmo Funding: \$108.4M
 OpenSearch Amazon Web Services ★ 6,559 MCap: \$929.7B	 Pandora Qiniu Funding: \$396.9M	 parseable Parseable ★ 954	 日志易 rizhiyi.com Rizhiyi Funding: \$11.4M	 sematext Sematext	 splunk Splunk MCap: \$14.8B
 sumo logic Sumo Logic MCap: \$1.4B	 Tencent Cloud Log Service Tencent Funding: \$404.3B	 TRINK Trink.io Trink.io			

ELK stack




- One common platform for logging is ELK: Elastic Search, Logstash and Kibana
- Logstash for log gathering and store
- Elastic Search to provide a query layer on top of logs
- Kibana a visualization platform

Fluentd



- Fluentd is an open source data collector for unified logging layer
- Is getting more space on logging inside a Kubernetes cluster
- Less resource consumption compared with ELK
- More integrated with components outside the cluster

Prometheus & Grafana



Motivation



- Kubernetes grant basic metrics about pods (memory, cpu)
- Those metrics are not sufficient when you want to have a better monitorization from your cluster and your applications
- Not only you need to gather new metrics, but you need a better visualization for them
- Dynamic and sharable dashboards are crucial for an efficient and proactive monitorization of any system and infrastructure

Prometheus



- Prometheus an open-source systems monitoring and alerting toolkit originally built at SoundCloud.
- It is now a standalone open source project and maintained independently of any company.
- Prometheus joined the Cloud Native Computing Foundation in 2016 as the second hosted project, after Kubernetes.
- Prometheus collects and stores its metrics as time series data, i.e. metrics information is stored with the timestamp at which it was recorded, alongside optional key-value pairs called labels.

Prometheus Features



- Prometheus's main features are:
 - a multi-dimensional data model with time series data identified by metric name and key/value pairs
 - PromQL, a flexible query language to leverage this dimensionality
 - no reliance on distributed storage; single server nodes are autonomous
 - time series collection happens via a pull model over HTTP
 - pushing time series is supported via an intermediary gateway
 - targets are discovered via service discovery or static configuration
 - multiple modes of graphing and dashboarding support

Prometheus PromQL



- Prometheus provides a functional query language called PromQL (Prometheus Query Language) that lets the user select and aggregate time series data in real time.
- The result of an expression can either be shown as a graph, viewed as tabular data in Prometheus's expression browser, or consumed by external systems via the HTTP API

```
# Latest sample
metric_name

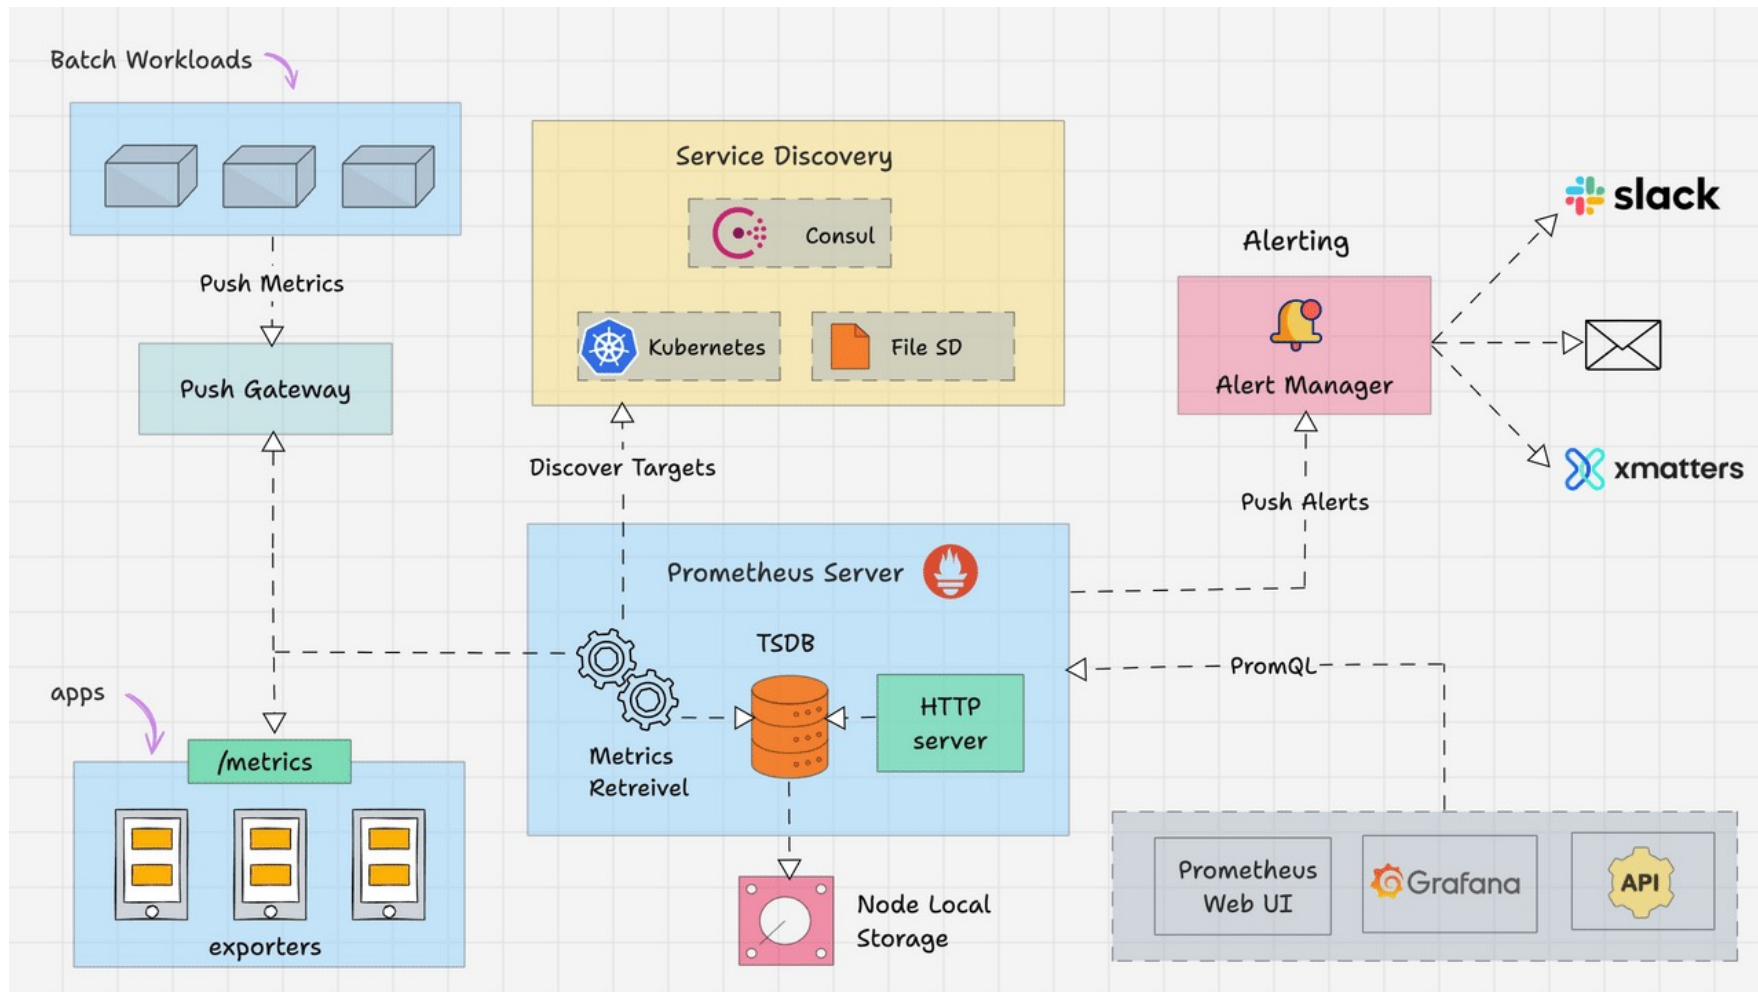
# Range
metric_name[5m]

# Labels
metric_name{label1="a",label2="b"}

# Functions
rate(metric_name[5m])
sum(metric_name)
delta(metric_name[5m])

# Comparisons
metric_name > 10*1024
```



Prometheus Architecture



Prometheus Components



Prometheus server
which scrapes and
stores time series data



Client libraries for
instrumenting
application code



Push gateway for
supporting short-lived
jobs



Alertmanager to
handle alerts



Visualization tools are
external but integration
with Grafana is natural



kubernetes

Grafana



- Grafana allows you to query, visualize, alert on and understand your metrics no matter where they are stored.
- Allow you to unify your data from several sources and make interactive dashboards
- Have a great linkage with Prometheus using PromQL to create dashboards
- Dashboards are described JSON what made really easy to share between the community

Demo | Prometheus & Grafana



Autoscaling



Motivation



- Kubernetes can handle several replicas of the same pods
 - ReplicaSets handle replication
 - Services handle load balancing between them
- However, if the demand of a service starts to grow, the number of replicas deployed may be not sufficient to handle requests
- Number of replicas can be changed manually but it's not scalable
- Kubernetes have a HorizontalPodAutoscaler (HPA) object to handle scalability of a Deployment automatically

Horizontal Pod Autoscaler



- Horizontal scaling means that the response to increased load is to deploy more Pods
- HPA defines a minimum and maximum number of replicas
- If the load increases, and the number of Pods is below the configured maximum, the HPA instructs the Deployment to scale up
- If the load decreases, and the number of Pods is above the configured minimum, the HPA instructs the workload resource to scale down
- HPA uses a control loop with a defined interval (default is 15 seconds) to check if some change is needed

Horizontal Pod Autoscaler



- To make the decision about scaling, HPA uses metrics about pods resources (CPU, Memory) utilization
- Metric target can be set as a percentage or raw value (preferable)
- Percentage value: controller calculates the utilization value as a percentage of the equivalent resource request
- Raw value: metric values are used directly

Horizontal Pod Autoscaler



- HPA uses a mean of the utilization or the raw value across all targeted Pods, and produces a ratio used to scale the number of desired replicas.
- Algorithm uses the following formula

$$\text{desiredReplicas} = \text{ceil}[\text{currentReplicas} * (\text{currentMetricValue} / \text{desiredMetricValue})]$$

- If metrics cannot be gathered from one pod, is considered as using 0% for scale up and 100% for scale down

HPA Metrics



- HPA can use 3 types of metrics to make the decision to scale up/down: per-pod resource metrics, custom metrics and external metrics
- Per-pod resource metrics: metrics gathered by native Metrics Server, like CPU, Memory and, GPU (near future)
- Custom metrics: metrics gathered by metrics scrappers plugin installed on the cluster, like Prometheus. For instance, you can autoscale your pods based on number of requests.
- External metrics: metrics that can be gathered from external resources using an additional plugin like Prometheus. For instance, you can autoscale your pods based on queued messages on a message queue.

HPA with Multiple Metrics



- You may specify more than one metric to be analyzed by HPA
- HPA makes the calculation for each metric and then select the biggest replica number from those calculations
- HPA can use multiple metrics from different sources

Other types of autoscaling



- [Vertical Pod Autoscaler](#): adjusts the resource requests and limits of a container
- [Cluster Autoscaler](#): adjusts the number of nodes of a cluster
- Both tools are defined and maintained by Kubernetes community but not available on a vanilla Kubernetes cluster
- VPA is fully implemented by Kubernetes community code
- Cluster Autoscaler needs to have specific implementation from nodes provider (and really hard to implement in on-prem cluster with bare metal 😊)

Questions?



kubernetes